

OM JAY MISHRA

Data Engineer · Azure Databricks Specialist

+91-9711107198 · mishra.omjay@gmail.com
[linkedin.com/in/omjay-mishra](https://www.linkedin.com/in/omjay-mishra) · Delhi NCR, India

SUMMARY

Data Engineer with 3+ years building scalable Lakehouse architectures on Azure Databricks. Reduced pipeline runtime from 6 hours to 30 minutes on 500M+ record datasets using Liquid Clustering, improved query performance 65% across 50M+ records, and cut Databricks DBU costs 25% through cluster optimization. Databricks Certified Data Engineer Associate.

TECHNICAL SKILLS

Databricks & Lakehouse	Azure Databricks, Delta Lake, Unity Catalog, Liquid Clustering, DBU & Cluster Optimization, Photon Engine, Medallion Architecture, PySpark
Azure Data Engineering	Azure Data Factory (ADF), ADLS Gen2, Azure DevOps, CI/CD Pipelines, Docker & Azure Containers
Programming & Querying	Python (PySpark, Pandas), SQL, Hive, REST APIs
Modeling & Governance	Star Schema Modeling, Data Governance, Data Quality Assurance, Metadata Management, Workflow Automation
BI & Visualization	Power BI, Tableau, Looker Studio

WORK EXPERIENCE

Data Engineer · ExaQube Technologies, Chennai Jun 2026 – Present

- Migrated 40+ production scripts from a legacy Ubuntu server to Databricks notebooks and Workflows, decommissioning the legacy environment
- Reduced a thrice-daily pipeline's runtime from ~1 hr 50 min to ~30 min per run by resolving data skew and migrating to job compute clusters, cutting daily compute hours by ~70% and lowering DBU costs
- Containerized a VPN-dependent pipeline using Docker on Azure containers; accelerated delivery through AI-assisted development (Claude + VS Code + Databricks)

Data Engineer · Azimuth GRC, Gurgaon May 2023 – May 2026

- Led end-to-end migration from legacy databases to Databricks Lakehouse, transforming SQL-based workflows into scalable PySpark and Delta Lake architectures
- Resolved partition skew on 500M+ record pipelines using Liquid Clustering, reducing runtime from 6 hours to 30 minutes
- Developed a Gold-layer Star Schema in Databricks, improving query performance by 65% across 50M+ records
- Reduced Databricks DBU consumption by 25% through cluster optimization, autoscaling, and Photon Engine utilization
- Orchestrated Azure Data Factory and Databricks pipelines with custom Python automation, eliminating 270+ hours of manual effort

Data Analytics Engineer · PlanetSpark, Gurgaon Nov 2022 – May 2023

- Automated data transformation workflows for 10K+ student records using Google Apps Script, improving data preparation efficiency
- Consolidated multi-source data into a centralized Power BI reporting layer, enabling stakeholders to track engagement and enrolment metrics

CERTIFICATIONS

- [Databricks Certified Data Engineer Associate](#)
- [Microsoft Azure Data Fundamentals \(DP-900\)](#)
- [Google AI Professional Certificate — Coursera](#)

EDUCATION

B.Tech, Mechanical & Automation Engineering
Maharaja Agrasen Institute of Technology, Delhi · 2019–2023
Languages: English, Hindi