

OM JAY MISHRA

Data Engineer · Azure Databricks Specialist

+91-9711107198 · mishra.omjay@gmail.com

[linkedin.com/in/omjay-mishra](https://www.linkedin.com/in/omjay-mishra) · Delhi NCR, India

Portfolio: omjay.github.io/Omjay-Portfolio

SUMMARY

Data Engineer with 3+ years designing and modernizing production Lakehouse platforms on Azure Databricks across ingestion, PySpark transformation, Delta Lake modeling, orchestration, data quality, governance, CI/CD, and cost-performance optimization. Reduced 500M+ record pipelines from 6 hours to 30 minutes, improved query performance 65%, and cut Databricks DBU consumption 25%.

TECHNICAL SKILLS

Databricks & Lakehouse	Azure Databricks, PySpark, Spark SQL, Delta Lake, Unity Catalog, Liquid Clustering, Photon, Medallion Architecture, Databricks Workflows
Azure Data Engineering	Azure Data Factory (ADF), ADLS Gen2, Azure DevOps, CI/CD, Docker, Azure Monitor, Job Clusters, Autoscaling
Programming & Querying	Python, SQL, Pandas, Hive, REST APIs, ETL/ELT, Incremental Loading, Metadata-Driven Pipelines
Modeling & Governance	Star Schema, Data Modeling, Data Quality, Data Governance, Audit Logging, Schema Evolution, Workflow Orchestration
BI & Visualization	Power BI, Tableau, Looker Studio

WORK EXPERIENCE

Data Engineer · ExaQube Technologies, Chennai Jun 2026 - Present

- Migrated 40+ production Python scripts from a legacy Ubuntu server to Databricks Workflows, replacing cron execution with dependency-aware orchestration, retries, alerting, and job clusters
- Reduced a thrice-daily pipeline's runtime from ~1 hr 50 min to ~30 min per run by resolving data skew and migrating to job compute clusters, cutting daily compute hours by ~70% and lowering DBU costs
- Built config-driven, Dockerized ingestion with idempotent checkpoints and schema evolution into Bronze Delta tables, allowing new sources to onboard without new pipeline code
- Tuned shuffle partitions, broadcast joins, adaptive query execution, file compaction, and clustering to stabilize heavier Spark workloads as data volume grew

Data Engineer · Azimuth GRC, Gurgaon May 2023 - May 2026

- Led migration from legacy databases and SQL workflows to an Azure Databricks Lakehouse, redesigning production pipelines in PySpark and Delta Lake with incremental loading patterns
- Reduced 500M+ record pipelines from 6 hours to 30 minutes using Liquid Clustering; improved Gold-layer query performance 65% across 50M+ records with a Star Schema
- Cut Databricks DBU consumption 25% through job clusters, autoscaling, Photon, and workload tuning; file compaction and Delta vacuuming reduced storage footprint 30%
- Automated ADF, Databricks, SQL, and Python workflows, eliminating 270+ hours of recurring manual work; metadata-driven ADF patterns reduced maintenance overhead 50%
- Implemented Unity Catalog access controls and audit logs, a PySpark Expectations framework for Bronze-to-Silver quality checks, and multi-stage Azure DevOps CI/CD

Data Analytics Engineer · PlanetSpark, Gurgaon Nov 2022 - May 2023

- Automated transformation and validation for 10K+ student records using Google Apps Script and Excel, cutting manual reporting cycles 40%
- Built 10+ engagement and retention dashboards; analysis informed product updates associated with 15% higher engagement and 20% better satisfaction metrics

CERTIFICATIONS

- [Databricks Certified Data Engineer Professional](#)
- [Databricks Certified Data Engineer Associate](#)
- [Microsoft Azure Data Fundamentals \(DP-900\)](#)
- [Google AI Professional Certificate](#)

EDUCATION

B.Tech, Mechanical & Automation Engineering

Maharaja Agrasen Institute of Technology, Delhi · 2019-2023

Languages: English, Hindi